

# N-body memory layout exploration


Bernhard Manfred Gruber  
CERN, CASUS

# N-body simulations

- Used to simulate interaction of many bodies
  - Galaxies, stars, atoms, electrons, proteins, ...
  - Gravity, electromagnetic forces, van der Waals forces, ...
- [https://www.youtube.com/results?search\\_query=nbody+simulation](https://www.youtube.com/results?search_query=nbody+simulation)
- $N^2$  computational complexity
  - There are faster methods and acceleration data structures
    - E.g. using Octrees and approx. groups of bodies
    - But they still use the  $N^2$  base algorithm inside

# Objective

- Given 3 n-body CPU implementations
  - ~430 LOCs C++
  - with different memory layouts
    - AoS (Array of Structs)
    - SoA (Structs of Arrays)
    - AoSoA (Array of Structs of Arrays)
- Find the fastest possible CUDA implementation
  - by finding the fastest memory layout for global and shared memory
  - without functional changes



# Tasks

1. Start by cloning: [https://github.com/bernhardmgruber/nbody\\_memory\\_layout\\_student\\_project](https://github.com/bernhardmgruber/nbody_memory_layout_student_project)
2. Compile/run sample CPU code. Compare runtime of different memory layouts
3. Inspect disassembly and understand CPU vectorization and access pattern: <https://godbolt.org/z/1jgo3z>
4. Write CUDA versions with various memory layouts in global memory and compare
5. Use a profiler to find explanations for the runtime differences
6. Write a faster CUDA variant using shared memory for caching
7. Find the best memory layout combination for global/shared memory